

Matrix Transposition of Big-Data

Sbonelo Mdluli:1101772, Heemal Ryan:792656, Haroon Rehman:1438756

ELEN4020A Project Due date: 17 May 2019

School of Electrical & Information Engineering, University of the Witwatersrand, Private Bag 3, 2050, Johannesburg, South Africa

Abstract—This project includes code written using MPI in order to solve a matrix transposition, this report shows cases and explores a performance comparison when using 16, 32 and 64 parallel processes. Using more process was found to be immediately faster however eventually became unnecessary as explained by Amdahls Law. Parallels I/O is used to access file simultaneously from different processes in order to maximize access speed.

I. INTRODUCTION

Big Data is often characterized by the 3 V's; Volume, extremely large data sets. Velocity, speed at which the data needs to be processed. Variety, referring to the different types of data that may need processing [1].

A very popular method of solving big data problems is parallel programming. Simply put, parallel programming allows for actions to be carried out simultaneously, in doing so, implementations are able to solve much larger problems than otherwise possible [2]. Efficient processing is usually required in numerical and scientific computations such as FFT and differential equations.

The following report focuses on a performance test using MPI (Message Passing Interface), a straight forward parallel programming method. The test entails recording the time taken to compute a matrix transposition. The following topics will be explored:

Problem Description What is MPI? Program Environment Code Structure Results Recommendations Conclusion

II. BACKGROUND

A. Problem Description

The problem entails showcasing a performance comparison when using 16, 32, 64 parallel processes to do a task. The task consists of performing a matrix transposition using MPI. The transposition should be done for a several sized matrices (see Tables 4-6). The input data is provided as per Tables 1 and 2

Table 1: Matrix A[8][8]

$a_{0,0}$	$a_{0,1}$	$a_{0,2}$	$a_{0,3}$	$a_{0,4}$	$a_{0,5}$	$a_{0,6}$	$a_{0,7}$
$a_{1,0}$	$a_{1,1}$	$a_{1,2}$	$a_{1,3}$	$a_{1,4}$	$a_{1,5}$	$a_{1,6}$	$a_{1,7}$
$a_{2,0}$	$a_{2,1}$	$a_{2,2}$	$a_{2,3}$	$a_{2,4}$	$a_{2,5}$	$a_{2,6}$	$a_{2,7}$
$a_{3,0}$	$a_{3,1}$	$a_{3,2}$	$a_{3,3}$	$a_{3,4}$	$a_{3,5}$	$a_{3,6}$	$a_{3,7}$
$a_{4,0}$	$a_{4,1}$	$a_{4,2}$	$a_{4,3}$	$a_{4,4}$	$a_{4,5}$	$a_{4,6}$	$a_{4,7}$
$a_{5,0}$	$a_{5,1}$	$a_{5,2}$	$a_{5,3}$	$a_{5,4}$	$a_{5,5}$	$a_{5,6}$	$a_{5,7}$
$a_{6,0}$	$a_{6,1}$	$a_{6,2}$	$a_{6,3}$	$a_{6,4}$	$a_{6,5}$	$a_{6,6}$	$a_{6,7}$
$a_{7,0}$	$a_{7,1}$	$a_{7,2}$	$a_{7,3}$	$a_{7,4}$	$a_{7,5}$	$a_{7,6}$	$a_{7,7}$

Table 2: Data *matrixFile.N*.

8	$a_{0,0}$	$a_{0,1}$	$a_{0,2}$	$a_{0,3}$	$a_{0,4}$	$a_{0,5}$	$a_{0,6}$	$a_{0,7}$	$a_{1,0}$	$a_{1,1}$	$a_{1,2}$...	$a_{7,6}$	$a_{7,7}$
---	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----	-----------	-----------

III. MESSAGE PASSING INTERFACE

MPI is a standardized message passing structure used in parallel computing (message passing refers to sending an instruction to a process such as an object, function or thread). MPI provides parallel hardware vendors with a set of clearly defined routines/instructions that may be efficiently implemented. Table 3 contains a summary of MPI's strengths and weaknesses [3], [4], [5].

Table 3: Strengths vs Weaknesses

Strengths	Weaknesses
Highly efficient	No support for multi-threading
Scalability	Inefficient dynamic process spawning
Easily available	No debugging facilities
Large volume of resources	
Reliable	
Stable communication	

IV. PROGRAM ENVIRONMENT

Code is to be tested using hornet01.eie.wits.ac.za, IP Address: 146.141.116.172, the host contains 16GB of available memory and a 3.4GHz CPU. A maximum of 8 threads is allowed per host. The environment is running on MPICH3.3 which is an MPI implementation. The maximum number of processors available on hornet01 is 373.

V. CODE STRUCTURE

Before exploring the programs code structure, we will explain how an MPI application works: MPI programs consist of multiple processes running at the same time. These processes have a unique identifiers(ranks) starting from 0 to (No. of Processes - 1) which run asynchronously. MPI also contain the notion of 'communicators', communicators describe a collection of processes, at initialization MPI_COMM_WORLD is made and contains all processes that were started upon running of the application [6].

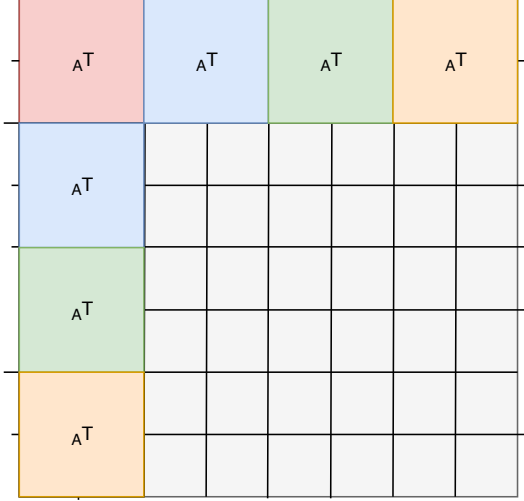


Fig. 1: Transposition Algorithm

The transposition algorithm is performed internally in blocks of 4 and those blocks are they put in to the overall encompassing matrix for further transposition. The final transposition is done by swapping the blocks. The blocks are put in a continuous rowtype where they are scattered to different processors, alltoall is used to perform the transpose of the bigger blocks.

The file generation and reading is done using parallel I/O which is one technique used to access data simultaneously from different application processes to maximize bandwidth and speed access time. This uses collective I/O which is used to move data between processors.

VI. RESULTS

The following results were found during testing of code writing using the hornet01.eie.wits.ac.za host. Appendix B contain line graph plots for each performance test. The timing includes processing and communication time, these are fairly similar with most of the inputs, so only the processing time is reported.

Table 4

No. of Processes = 16			
Array Size	Read(ms)	Write(ms)	Transpose(sec)
2^3	74.59	0.119	0.930
2^4	70.02	0.294	0.080
2^5	86.54	3.360	0.079
2^6	96.50	23.98	0.101
2^7	112.9	26.38	0.094

Table 5

No. of Processes = 32			
Array Size	Read(ms)	Write(ms)	Transpose(sec)
2^3	70.55	0.129	0.177
2^4	77.21	0.208	0.213
2^5	80.49	1.567	0.221
2^6	83.41	11.08	0.154
2^7	95.32	10.39	0.210

Table 6

No. of Processes = 64			
Array Size	Read(ms)	Write(ms)	Transpose(sec)
2^3	66.77	0.093	0.293
2^4	67.61	0.120	0.321
2^5	70.56	0.255	0.329
2^6	72.16	2.782	ONCLUSION0.329
2^7	95.32	2.752	0.285

The results are according to Amdahl's law which can be used to estimate performance. The principle states that increasing the number of processors does not necessarily increase speed up time.

VII. RECOMMENDATIONS

The current transposition algorithm waits for rank 0 to do the internal block matrix locally and sends the results the different processors. This is inefficient because this causes a block/barrier. To increase the computation time the block transpose should be done in other processors as not to cause a barrier. Further optimizations can be made by making use of the shared memory of each distributed memory. This can be done by using Pthreads in each node. The send can be made asynchronous and non-blocking.

VIII. CONCLUSION

This report presents an investigation of matrix transpose and parallel I/O across multiple processors. A $N \times N$ square matrix of $N = 2^n$ is where $N = 3, 4, 5, 6, 7$ is transposed using block transposition and MPI_alltoall. The input data is generated into a binary file using parallel I/O. The results from the investigation are consistent with theoretical studies namely Amdahl's law.

REFERENCES

- [1] SearchDataManagement. (2019). What is big data? - Definition from WhatIs.com. [online] Available at: <https://searchdatamanagement.techtarget.com/definition/big-data> [Accessed 16 May 2019].
- [2] O'Reilly — Safari. (2019). Patterns for Parallel Software Design. [online] Available at: https://www.oreilly.com/library/view/patterns-for-parallel/9780470697344/9780470697344_advantages_and_disadvantages_of_paral.html [Accessed 16 May 2019].
- [3] GeeksforGeeks. (2019). MPI - Distributed Computing made easy - GeeksforGeeks. [online] Available at: <https://www.geeksforgeeks.org/mpi-distributed-computing-made-easy/> [Accessed 16 May 2019].
- [4] Cecs.wright.edu. (2019). [online] Available at: <http://cecs.wright.edu/~schung/ceg820/mpi-book.pdf> [Accessed 16 May 2019].
- [5] Rosul, C. (2019). Message Passing Interface (MPI) Advantages and Disadvantages for applicability in the NoC Environment by. [online] SemanticScholar.org. Available at: [https://www.semanticscholar.org/paper/Message-Passing-Interface-\(-MPI-\)-Advantages-and-in-Rosul/6c5dd34a693f72adaebd1d02d521c5b7cfd589](https://www.semanticscholar.org/paper/Message-Passing-Interface-(-MPI-)-Advantages-and-in-Rosul/6c5dd34a693f72adaebd1d02d521c5b7cfd589) [Accessed 16 May 2019].
- [6] Userinfo.surfsara.nl. (2019). MPI in two hours — userinfo.surfsara.nl. [online] Available at: <https://userinfo.surfsara.nl/systems/shared/mpi/mpi-intro> [Accessed 16 May 2019].

APPENDIX

A. Pseudo-code

Algorithm 1: Function that reads input data

Function name: *Main Function*
Input/Output Files

char *in = argv [1], char *out = argv[2],

Setup

int rank, numProcs, inputError, SIZE,

int *sizeChecker, *tempBuffer, n

Initialize MPI

MPI_Init(argc, argv)

MPI_Comm_rank(MPI_COMM_WORLD, rank)

MP_Comm_size(MPI_COMM_WORLD, numProcs)

MPI setup

MPI_File fh

MPI_Win win

MPI_Status status

MPI_Offset disp

MPI_Offset blockSize

Open File

inputError = MPI_File_open(MPI_COMM_WORLD, in, MPI_MODE_RDONLY,MPI_INFO_NULL,fh)

if *There is an input error = true* **then**

| exit;

else
Get size of Matrix

sizeChecker = (int*)malloc(sizeof(int))

MPI_File_read(fh, sizeChecker, 1, MPI_INT, status)

SIZE = sizeChecker[0]

n = (int)(SIZE/numProcs)

Condition the buffer for each rank blockSize = SIZE*n

disp = ((rank*blockSize)+1)*sizeof(int)

tempBuffer = (int*)malloc(SIZE*n*sizeof(int))

Read the file into the buffer MPI_File_set_view(fh, disp, MPI_INT,
MPI_INT,"native",MPI_INFO_NULL)

MPI_File_read(fh,tempBuffer,blockSize,MPI_INT, status)

MPI_File_close(fh)

end

return 0

Algorithm 2: Function performs matrix transpose

Function name: *Main Function*
Initialize

```

int i, world_rank, world_size
double process_time, process_time_to_avg, comm_time = 0
int **rand_matrix, **trans_matrix = NULL
int *recv_buffer = (int*) malloc(mat_size * sizeof(int))
int *recv_buffer2 = (int*) malloc(mat_size * sizeof(int))
MPI_Init(NULL, NULL)
MPI_Comm_rank(MPI_COMM_WORLD, world_rank)
MPI_Comm_size(MPI_COMM_WORLD, world_size)

```

Scatter rows

```

if world_rank == 0 then
    MPI_Scatter((rand_matrix[0][0]), mat_size, MPI_INT,
    (recv_buffer[0]), mat_size, MPI_INT, 0, MPI_COMM_WORLD)
else
    Send different rows
    MPI_Scatter(NULL, 0, MPI_INT, (recv_buffer[0]), mat_size, MPI_INT, 0, MPI_COMM_WORLD)
end

```

Do transpose of scattered rows MPI_Alltoall(recv_buffer, 1, MPI_INT, recv_buffer2, 1, MPI_INT, MPI_COMM_WORLD)

Gather to form matrix if world_rank == 0 then

```

    comm_time += MPI_Wtime()
    trans_matrix = create_matrix(mat_size, mat_size)
    comm_time -= MPI_Wtime()
    MPI_Gather(recv_buffer2, mat_size, MPI_INT, (trans_matrix[0][0]), mat_size, MPI_INT, 0,
    MPI_COMM_WORLD)
else
    MPI_Gather(recv_buffer2, mat_size, MPI_INT, NULL, 0, MPI_INT, 0, MPI_COMM_WORLD)
end
return 0

```

Algorithm 3: Write to output file

Function name: *Main Function*
Initialize

```

int i, j, rank, size, N=128
srand(time(0))
MPI_File fhw
MPI_Status status
MPI_Offset offset
MPI_Offset chunk
MPI_Init(argc, argv)
MPI_Comm_rank(MPI_COMM_WORLD, rank)
MPI_Comm_size(MPI_COMM_WORLD, size)
int buf[(N*N)/size]
int tmp

MPI_File_open(MPI_COMM_WORLD, "data", MPI_MODE_CREATE—MPI_MODE_WRONLY,
  MPI_INFO_NULL, fhw)
offset = rank*(N*N/size)*sizeof(int)
chunk = N*N/size
for  $i=0$ ;  $i$  less than  $(N*N)/size$  ;  $i++$  do
  int tmp = rand()%1000
  for  $j=0$ ;  $j$  less than  $i$ ;  $j++$  do
    if  $tmp = buf[j]$  then
      tmp = rand()
    else
      end
    buf[i] = tmp
  end

MPI_File_write_at_all(fhw, offset, buf, chunk, MPI_INT, status); MPI_File_close(fhw)
MPI_Finalize()
return 0

```

Algorithm 4: Void Transpose Function: Transposes all inner blocks of matrix

Function name: **void transpose-blocks*

Initialization :int thread-counter = 0

long local = thread-counter

thread-counter++, int pos = 0

int posY = 0 , int tmp = size-mat/block-len

int x = 0 , tensor block

while 1 **do**

 pos = (block-len*local)

 x = local/tmp

 posY = x*block-len

if posY \neq size-mat **then**

 break

else

 block = getSubMatrix(posY,pos)

 transpose(block)

 setSubMatrix(posY,pos,block)

 local= next-pos

 next-pos++

end

end

B. Performance Tests

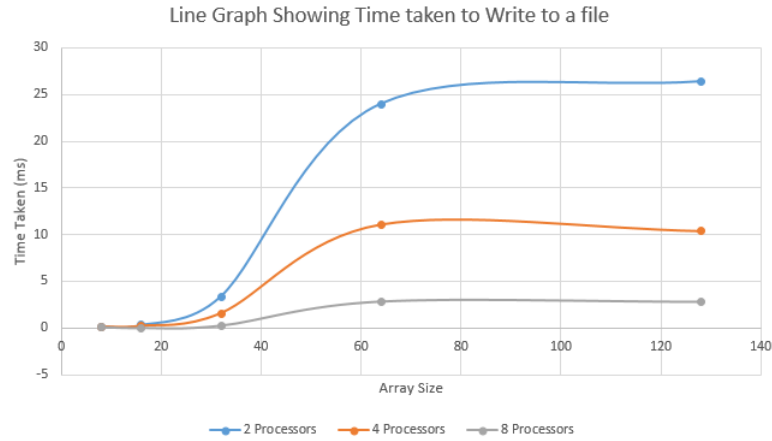


Figure 1: Time to write to File

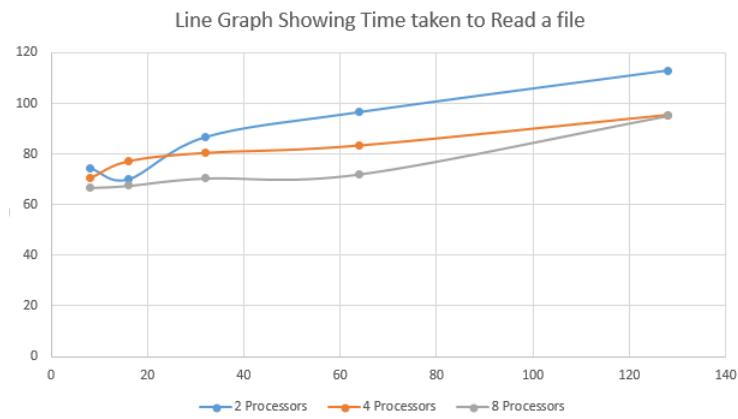


Figure 2: Time to Read from a file

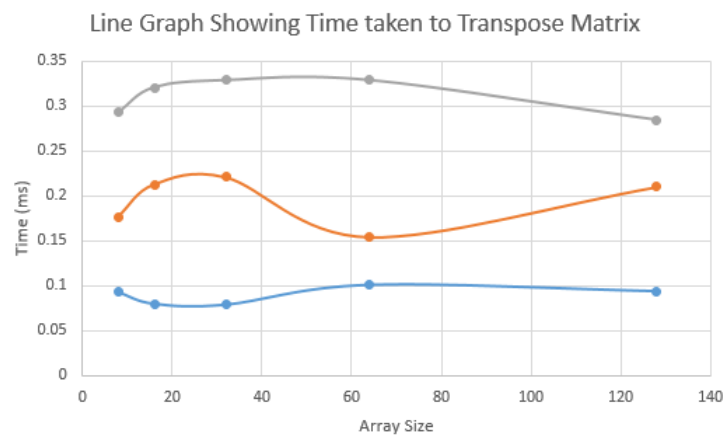


Figure 3: Time to Transpose Matrix